

CS4100 - Project 1 Report

Jansen Craft & Kaleb Demaline

What regular expressions/tokens did you identify, and how does your program react to each one.

Macros

F	[a-zA-Z0-9_]
D	[0-9]
INT	"int" "short" "unsigned int" "unsigned short" "long" "unsigned long" "long long" "unsigned long long"
FLOAT	"float" "double" "long double"

Regular Expressions and Reactions with Token names

Regular Expression	Reaction	Token Returned (If applicable)
<code>\\n</code>	Increment linecount	no token – single-line comment
<code>"/**([^\n]*[^\n])**/"</code>	Add number of newlines to linecount	no token – multi-line comment
<code>#include <[a-z.]*></code>	Do nothing	no token – include statement
<code>#include "[a-z.]*"</code>	Do nothing	no token – include statement
<code>\(</code>	Return token for open parenthesis	TK_OPEN_P
<code>\)</code>	Return token for closed parenthesis	TK_CLOSED_P
<code>\{</code>	Return token for open curly brace	TK_OPEN_CB
<code>\}</code>	Return token for closed curly brace	TK_CLOSED_CB
<code>\[</code>	Return token for open square bracket	TK_OPEN_SB
<code>\]</code>	Return token for closed square bracket	TK_CLOSED_SB
<code>;</code>	Return token for semicolon	TK_SEMICOLON
<code>,</code>	Return token for comma	TK_COMMA
<code>"[^\n]*"</code>	Add number of newlines to linecount. Return token for string literal	TK_STRING_LITERAL
<code>'[^\\]{1,2}'</code>	Return token for character literal	TK_CHAR_LITERAL
<code>([D]*\.[D]*) ([D])+</code>	Return token for number literal	TK_NUMBER_LITERAL
<code>\+</code>	Return token for plus	TK_PLUS
<code>\-</code>	Return token for minus	TK_MINUS
<code>!</code>	Return token for exclamation	TK_EXCLAMATION
<code>~</code>	Return token for tilde	TK_TILDA
<code>\&</code>	Return token for ampersand	TK_AMPERSAND
<code>*</code>	Return token for asterisk	TK_ASTERISK
<code>\=</code>	Return token for equal	TK_EQUAL

/	Return token for forward slash	TK_FORWARD_SLASH
\%	Return token for percent	TK_PERCENT
\^	Return token for caret	TK_CARET
>	Return token for greater	TK_GREATER
<	Return token for less	TK_LESS
\?	Return token for question	TK_QUESTION
\	Return token for pipe	TK_PIPE
{INT}	Return token for int	TK_INT
{FLOAT}	Return token for float	TK_FLOAT
"void"	Return token for void	TK_VOID
"char"	Return token for char	TK_CHAR
"while"	Return token for while	TK_WHILE
"for"	Return token for for	TK_FOR
"do"	Return token for do	TK_DO
"if"	Return token for if	TK_IF
"else"	Return token for else	TK_ELSE
"switch"	Return token for switch	TK_SWITCH
"case"	Return token for case	TK_CASE
"default"	Return token for default	TK_DEFAULT
"break"	Return token for break	TK_BREAK
"continue"	Return token for continue	TK_CONTINUE
"return"	Return token for return	TK_RETURN
"goto"	Return token for goto	TK_GOTO
{F}+	Return token for identifier	TK_IDENTIFIER
[\t\n\f]	Add number of newlines to linecount	no token - whitespace
.	Do nothing	no token - unmatched

In the main of the lex program, we print out the number associated with the token (1-42) with width 2. We prepend the number with zeros if it is < 10.

A brief description in your own words of how you implemented the Winnowing algorithm

```
int main(){
    int k = 20, w = 20;
    hash<string> hasher;
    vector<string> filenames;
    vector<vector<size_t>> fingerprints;

    string filename, tokens;
```

```

while(cin >> filename && getline(cin, tokens)){
    filenames.push_back(filename);
    auto end = remove_if(tokens.begin(),tokens.end(),[](char c){return c == ' '});
    tokens.erase(end,tokens.end());

    vector<string> k_grams;
    for (int i = 0; i <= tokens.size()-k; i++)
        k_grams.push_back(tokens.substr(i, k));

    auto fp = winnow(w,k_grams,hasher);
    fingerprints.push_back(fp);
}

```

In our main, we read in all the tokens of a file as one space-delimited string, and then remove all the whitespace. After this, we add the k_grams (substrings of this mega-string) to a vector. We have a hasher object from the stl library which maps strings to SIZE_T. We then pass the window size, the vector of k_grams, and the hasher object to our winnowing function.

```

vector<size_t> winnow(int w, const vector<string>& k_grams, hash<string>& hasher){
    deque<size_t> buffer(w, SIZE_T_MAX);
    vector<size_t> fingerprints;
    int min_hash_index = 0;

    // Load initial w-1 k grams
    for (int k_idx = 0; k_idx < w-1; k_idx++){
        buffer.push_back(hasher(k_grams[k_idx]));
        buffer.pop_front();
    }

    for (int k_idx = w; k_idx < k_grams.size(); k_idx++){
        buffer.push_back(hasher(k_grams[k_idx]));
        buffer.pop_front();
        min_hash_index--;

        if (min_hash_index == -1){
            min_hash_index = w-1;
            for(int i = w-1; i >= 0; i--){
                if (buffer[i] < buffer[min_hash_index]) min_hash_index = i;
            }
            fingerprints.push_back(buffer[min_hash_index]);
        } else {

```

```

        if (buffer.back() <= buffer[min_hash_index]){
            min_hash_index = w-1;
            fingerprints.push_back(buffer[min_hash_index]);
        }
    }
}
return fingerprints;
}

```

In the winnowing function, we initialize a deque to w nodes with `SIZE_T_MAX`. We then load in the first $w-1$ hashed k _grams into the deque by adding to the back and popping off the front. After this, starting at the w th hash, we slide add the next hash, remove the last one, increment our `min_hash_index` and then check to see if we have a new smallest hash. Firstly, if our best hash was just popped off the front of buffer, we find the rightmost smallest hash and set that as our minimum hash, making sure to record the hash and location to a fingerprints vector. If our best hash is still in the buffer, we just check it against the newest hash. If the new one is smaller, we set that our `min_hash_index` and then append its information to the fingerprints vector which is returned at the end of the function.

```

// Compare Fingerprints
for (int i = 0; i < fingerprints.size(); i++)
    sort(fingerprints[i].begin(), fingerprints[i].end());

vector<pair<float, string>> similarity_scores;

for (int i = 0; i < fingerprints.size(); i++){
    for (int j = i+1; j < fingerprints.size(); j++){
        vector<size_t> intersection;
        set_intersection(
            fingerprints[i].begin(), fingerprints[i].end(),
            fingerprints[j].begin(), fingerprints[j].end(),
            back_inserter(intersection)
        );
        similarity_scores.push_back({(float)intersection.size() / (float)fingerprints[i].size(), filenames[i] + ' ' + filenames[j]});
        intersection.clear();
    }
}

sort(similarity_scores.rbegin(), similarity_scores.rend());

for (auto score: similarity_scores){
    cout << setprecision(3) << setfill('0') << fixed << score.first << "t" << score.second << "\n";
}

```

}

In the main, once we have the fingerprints vectors for each of the files passed to standard input, we start our pairwise comparison. For the `set_intersection` from `stl` that we are using, we must first sort all the fingerprint vectors. After this, we can get the size of the intersection of fingerprint vectors of the files, divide that by the length of the file in the outer loop, and then use this as our percentage of common fingerprints found. This percentage is printed out for that comparison after the percentages are sorted at the end.

The results of your analysis, including identifying any submissions found by the algorithm that you believe may be plagiarism

1.000	bills_53.c bills_54.c
1.000	bills_09.c bills_54.c
1.000	bills_09.c bills_53.c
1.000	bills_03.c bills_54.c
1.000	bills_03.c bills_53.c
1.000	bills_03.c bills_09.c
0.971	bills_10.c bills_30.c
0.960	bills_47.c bills_48.c
0.904	bills_15.c bills_22.c

Here, at the top of our Plagiarism Report, we have the pairs of student submission that were most similar. The first thing that must be said is that this assignment was clearly very simple, and they were all given the same basic instructions on how to write the program. This makes it so most of the submissions are very similar by default. However, students 53, 54, 9, and 3 had the same fingerprints. This is so unlikely to happen naturally, that it should be safe to say that these students plagiarized. One could assign a cutoff of perhaps 90% similarity and could reasonably say these students cheated, but the safest bet is to at least talk to the students mentioned above.